

MP5

1. Overview

In `main.py` I implement a full Canny edge detector pipeline from scratch in pure NumPy:

- **Gaussian smoothing** with a user-defined kernel size $N \times N$ and standard deviation σ .
- **Gradient computation** $[G_x, G_y]$: $\text{magnitude} = \sqrt{G_x^2 + G_y^2}$ and $\arctan()$ for theta.
- **Automatic threshold selection**: histogram of the gradient magnitudes to CDF for picking T_{high} so that a given fraction of pixels are below it, and $T_{\text{low}} = 0.5 \times T_{\text{high}}$.
- **Non-maximum suppression** to thin edges, using either 4-way quantization or interpolation along the gradient.
- **Hysteresis thresholding & edge linking**: start from strong edges $\geq T_{\text{high}}$ and recursively follow any weak neighbors $\geq T_{\text{low}}$.
- **Parameter sweeps** over `N`, `sigma`, and `percentage_non_edge` to see how they affect final edge quality.

I test on four images (`lena.bmp`, `test1.bmp`, `pointer1.bmp`, `joy1.bmp`) and compared edge detection results with **OpenCV**.

2. Algorithm Description

1. Gaussian smoothing

- Build an $N \times N$ Gaussian kernel.
- Convolve it with the image (edge-pad) per channel.

2. Gradient computation

- Convert RGB to Grayscale.
- Compute gradients along X and Y and convert to magnitude and theta.

3. Threshold selection

- Flatten gradient magnitude into a histogram.
- Choose T_{high} at the bin where CDF is greater equal than `percentage_non_edge` and set $T_{\text{low}} = 0.5 \times T_{\text{high}}$

4. Non-maximum suppression

- For each pixel, compare its magnitude to the two neighbors along the gradient direction: pick one of 4 directions ($0^\circ, 45^\circ, 90^\circ, 135^\circ$) based on $[G_x, G_y]$ and signs.

5. Hysteresis & edge linking

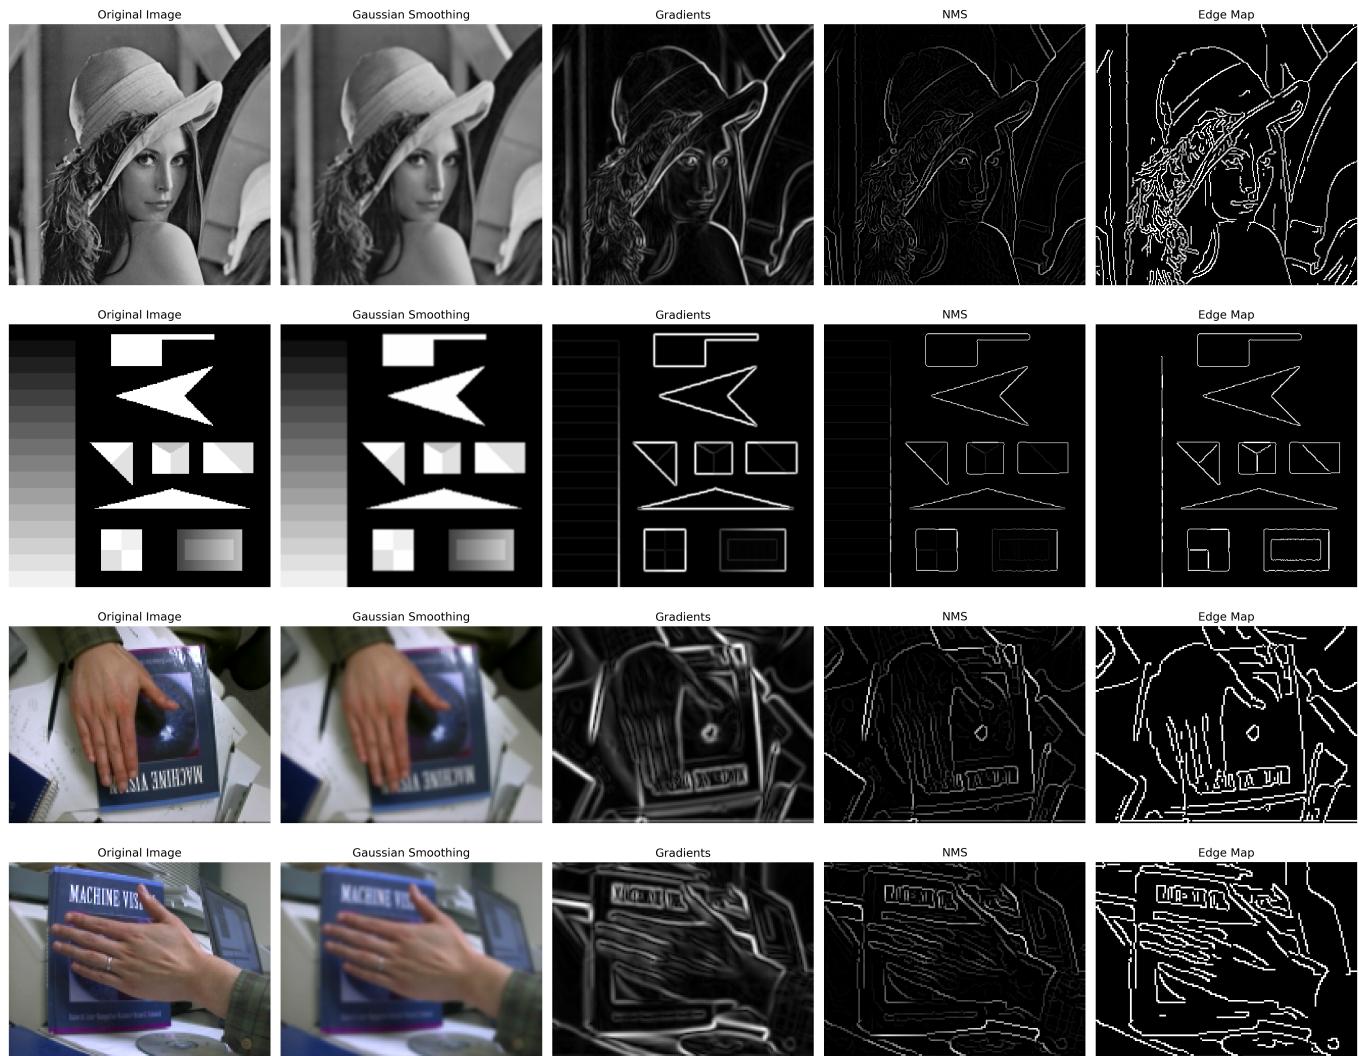
- Build boolean masks
- $\text{mag_low} = \text{mag_sup} \geq T_{\text{low}}, \text{mag_high} = \text{mag_sup} \geq T_{\text{high}}$
- For each unvisited strong pixel, do a stack-based flood fill over 8-neighbors in `mag_low`, marking all reached as final edges.

6. Parameter experiments

- Sweep three sets of parameter triples N , σ , $\text{percentage_non_edge}$ and plot the resulting edge maps side-by-side for each test image.

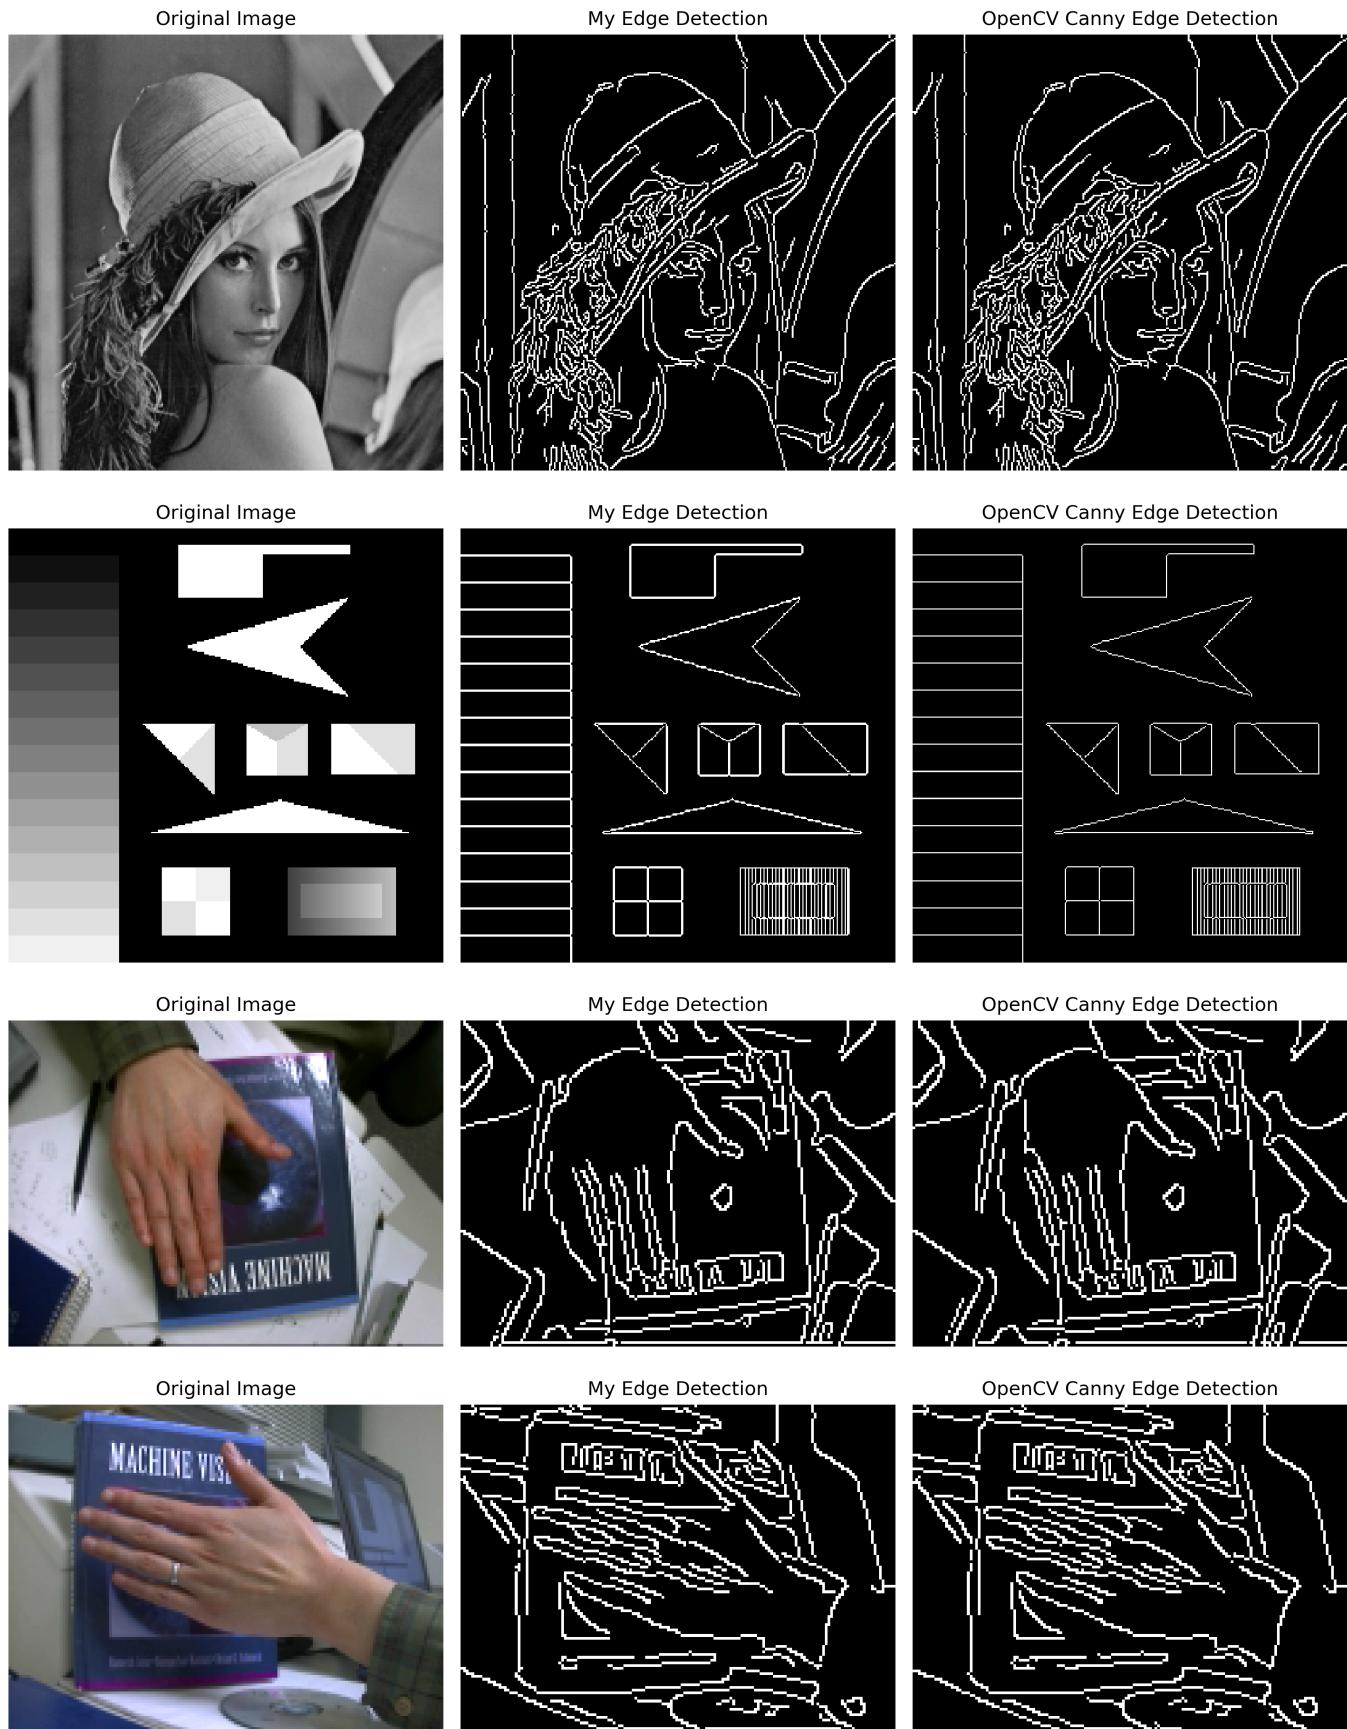
3. Results

Below are results about gaussian smoothing, gradient magnitudes, NMS and final edge maps. Results are showing all functions are working correctly.



3. Comparison

I compared the edge detection results with `cv2` functions. I found they're basically the same.

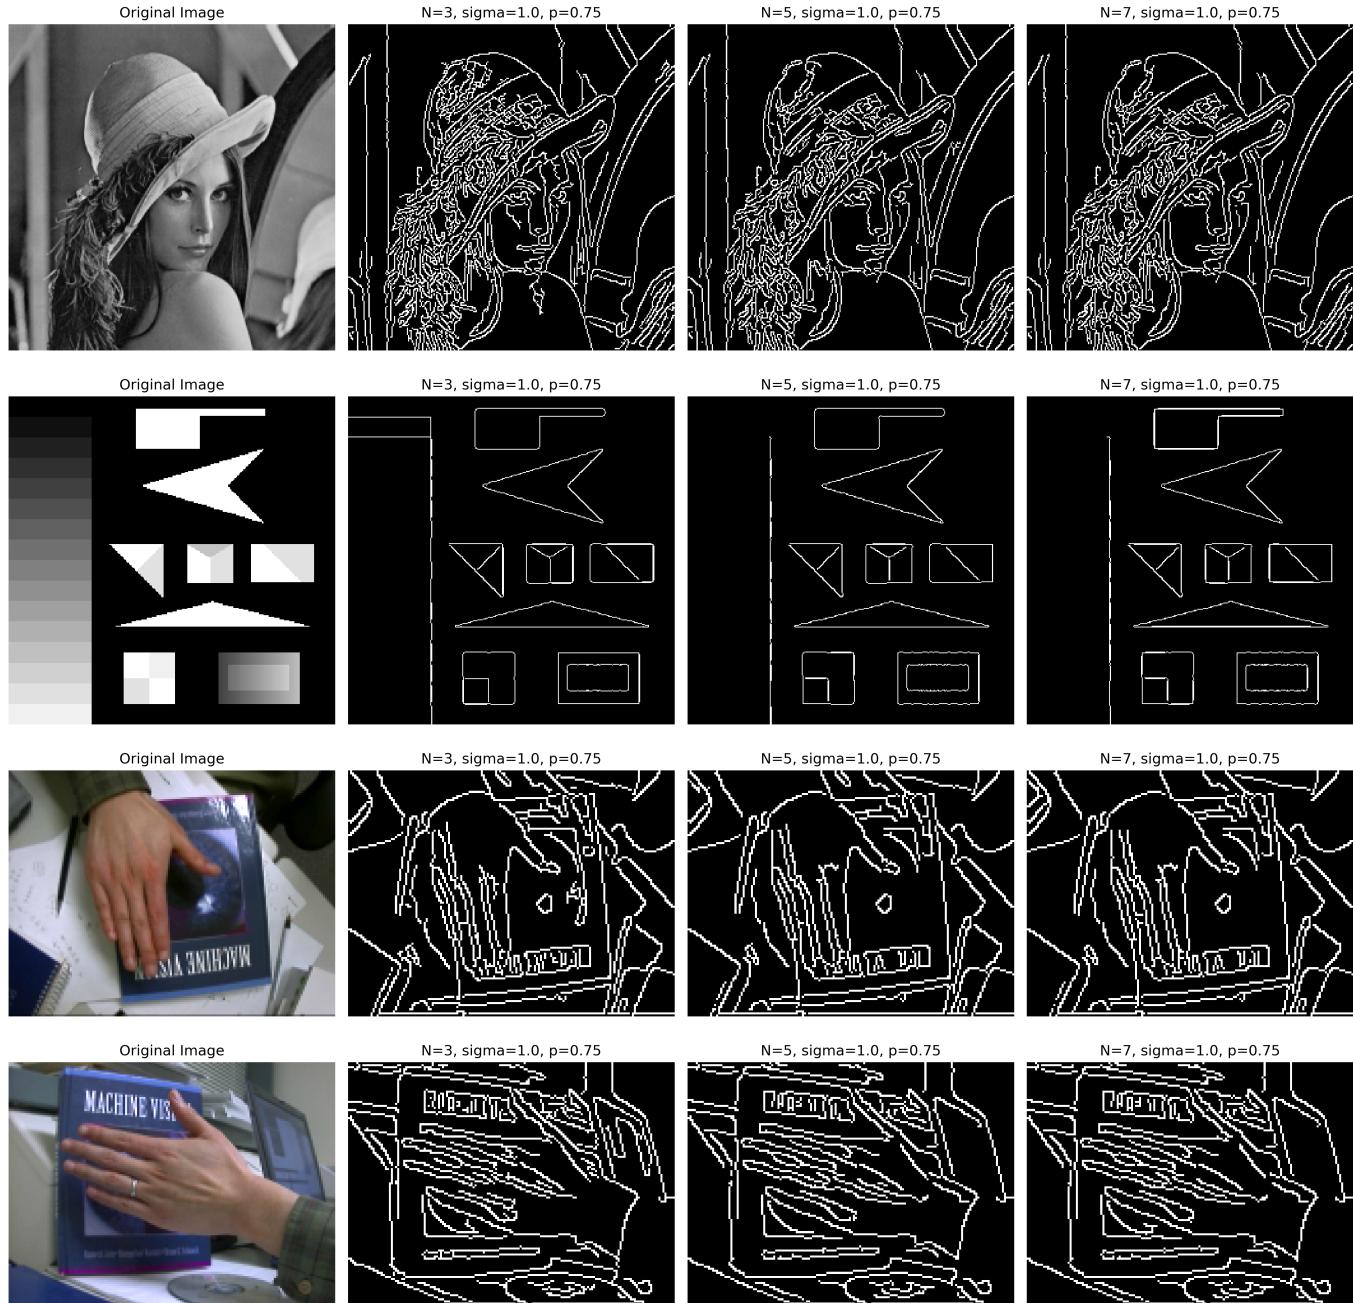


4. Tuning Parameters

I ran three experiments on each of the four test images (`lena`, `test1`, `pointer1`, `joy1`):

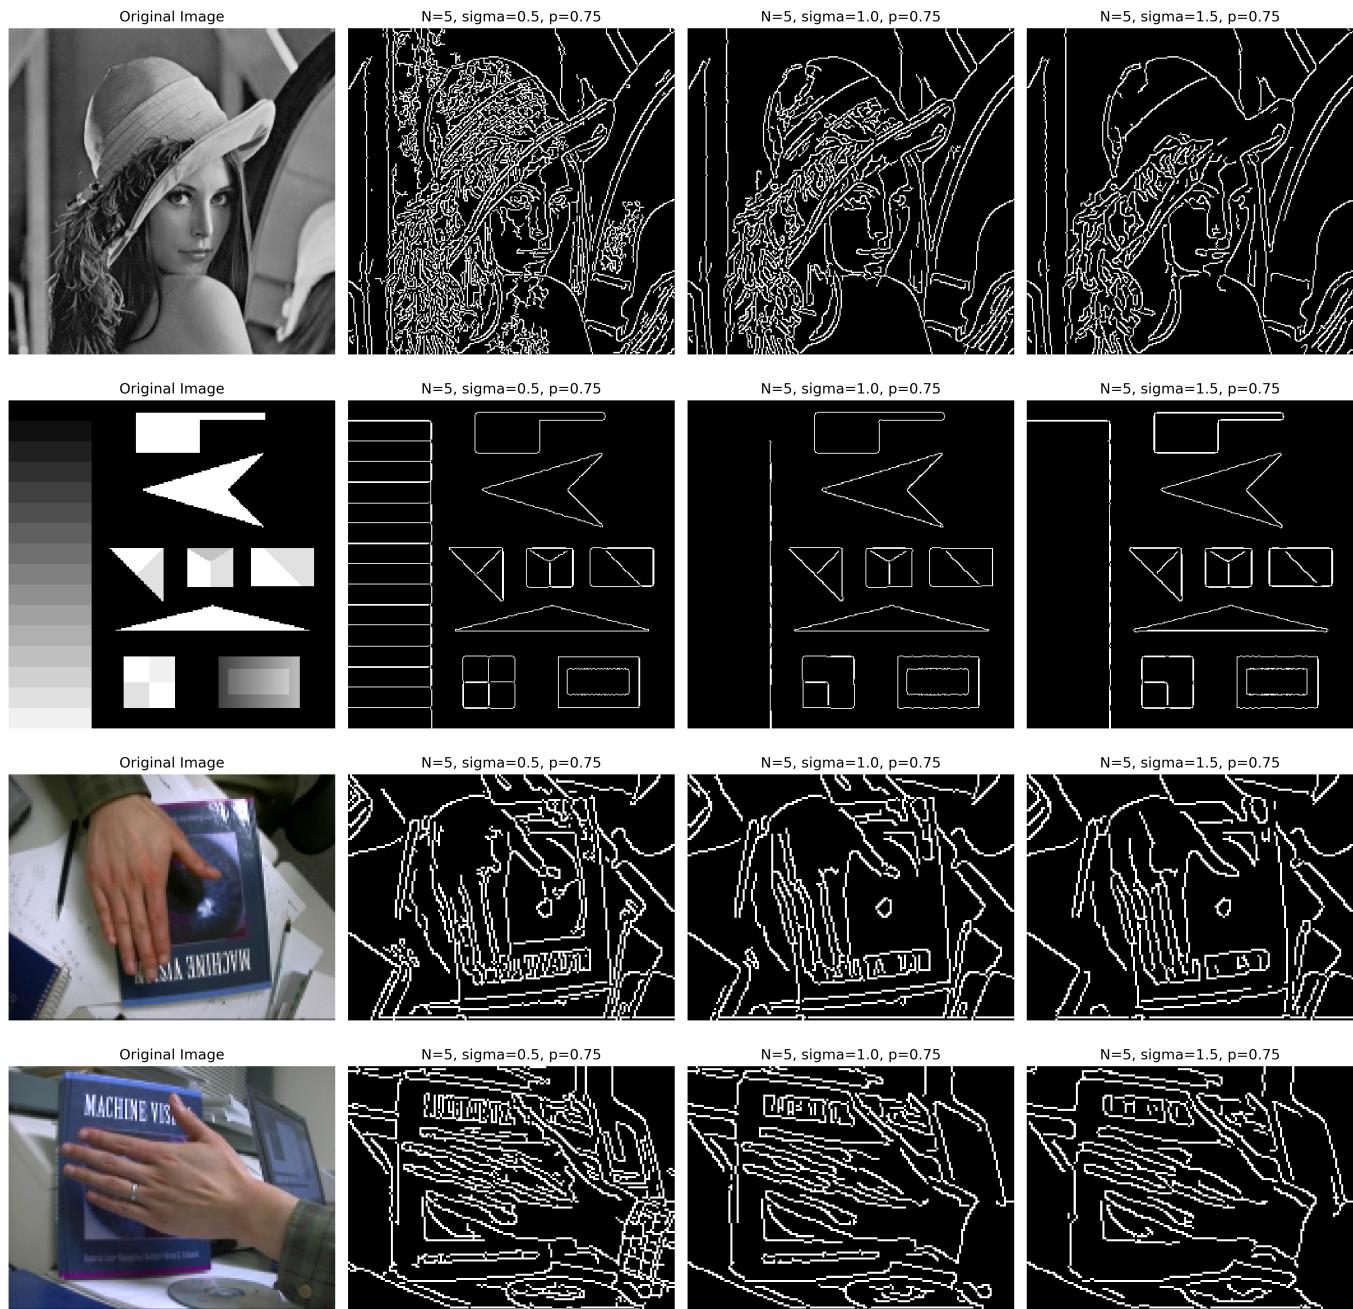
4.1. Varying Kernel Size N ($\sigma=1.0$, $p=0.75$)

I found there's no significant change when changing **N**.



4.2. Varying **sigma** (**N**=5, **p**=0.75)

I found the image will lose more details as **sigma** grows.



4.3. Varying Non-Edge Percentage `percentage_non_edge` ($N=5$, $\sigma=1.0$)

I found the edge would focus on more obvious contours as `percentage_non_edge` grows.

